

GPU accelerated Monte Carlo simulations of lattice spin models

M. Weigel, T. Yavors'kii

Institut für Physik, Johannes Gutenberg-Universität Mainz, Staudinger Weg 7, D-55099 Mainz, Germany

Abstract

We consider Monte Carlo simulations of classical spin models of statistical mechanics using the massively parallel architecture provided by graphics processing units (GPUs). We discuss simulations of models with discrete and continuous variables, and using an array of algorithms ranging from single-spin flip Metropolis updates over cluster algorithms to multicanonical and Wang-Landau techniques to judge the scope and limitations of GPU accelerated computation in this field. For most simulations discussed, we find significant speed-ups by two to three orders of magnitude as compared to single-threaded CPU implementations.

Keywords:

spin models; Monte Carlo simulations; GPU computing; cluster algorithms; generalized-ensemble simulations

1. Introduction

Due to their rich content in phase transitions and critical phenomena and the wealth of applications in condensed-matter physics and beyond, classical spin systems have been studied numerically in an enormous body of published research to date [1]. The ever increasing demand in computational power for computer simulations of spin models has prompted researchers interested in these problems to regularly use the cutting-edge computer technology of their time. In the past and present this has included the design and implementation of special-purpose computers for spin model simulations such as the cluster processor for studying critical phenomena in the ferromagnetic Ising model [2] or, more recently, the Janus machine for simulations of spin-glass models [3]. The construction of such machines entails an effort in financial and human resources significantly exceeding that of programming a regular computer or compute cluster. With the advent of general purpose computing on graphics processing units (GPGPU) [4] a high-performance computing architecture is gaining more and more widespread use in the scientific community [5] that is special-purpose in its original objective of displaying graphics and certainly less general in its applicability and flexibility than current CPUs but, nevertheless, significantly more easily accessible than the special-purpose machines mentioned above.

The possibility of harvesting the nominally vast computational power provided by current-generation GPUs rests on the degree to which a number of challenges can be met: (a) the efficient programming of such devices without the need of using (too) low-level techniques, (b) the suitability of the problem and the algorithms employed for calculations in a massively parallel setup, and (c) the degree to which peculiarities of such devices, such as special hierarchies of cache memories, difficulties with thread scheduling, or reduced precision in floating-point arithmetics, can be catered for in the implementations [6]. While a device independent programming interface has recently become available with the OpenCL framework [7], we here use the NVIDIA CUDA toolkit [8] due to its currently more mature status. It is the purpose of this contribution to summarize our experience with the suitability of a number of popular algorithms for the simulation of spin models for the massively parallel environment provided by current GPUs. In the

course of describing the results for various models and algorithms, we will discuss the question in how far architectural limitations and peculiarities represent significant hurdles for the efficient use of GPUs.

2. Metropolis algorithm

We studied classical $O(n)$ spin models with Hamiltonian

$$\mathcal{H} = - \sum_{\langle ij \rangle} J_{ij} \mathbf{s}_i \cdot \mathbf{s}_j, \quad (1)$$

where $n = 1$ corresponds to the discrete Ising model and $n = 2$ and $n = 3$ describe the continuous XY and Heisenberg models, respectively. The spins are located on square ($d = 2$) or simple cubic ($d = 3$) lattices and interact with nearest neighbors only. We first considered simulations using only single spin-flip moves accepted according to the Metropolis criterion [9],

$$p_{\text{acc}}(s_i \mapsto s'_i) = \min \left[1, e^{-\beta \Delta E} \right]. \quad (2)$$

We also considered updates according to the heat-bath prescription [1] which lead to only slightly (and locally) modified implementations and very similar speed-ups on GPUs as compared to CPU implementations and are, therefore, not discussed in detail here.

For an implementation on GPU, one needs to allow for the parallel update of a large number of spins. This is most straightforwardly accomplished for the case of nearest-neighbor interactions by making use of a checkerboard decomposition of the lattice [10]. To make efficient use of *shared memory* that can be accessed concurrently with very low latencies from all threads on a multiprocessor, we use a two-level hierarchical decomposition, cf. the left panel of Fig. 1 [11, 12]. All spins of one of the big tiles plus some boundary layer are collaboratively loaded into shared memory and subsequently updated by the individual threads of a thread block [6]. Inside of each tile, the checkerboard arrangement allows for all spins of one sub-lattice to be updated concurrently before a synchronization barrier occurs and the second sub-lattice is treated analogously. To amortize the effort of loading tiles into shared memory, a number k of updates of all spins of all big tiles of one color is performed before updating the second sub-lattice. Depending on the size of the tiles, this slows down the decorrelation of spin configurations. This effect, however, is more than counter-balanced by the performance increase, even close to criticality [12]. For good performance, a number of additional tricks are employed, including a pre-tabulation of the Boltzmann factors in Eq. (2) while storing this table as a texture [6], and generation of random numbers even if they are not required to reduce thread divergence. The simulation code can be downloaded at the authors' web site [13]. For the comparisons discussed here, an array of simple 32-bit linear congruential pseudo-random number generators (one per thread) is used. Although these are known to have rather poor properties, for the purpose at hand they appear to be sufficient even for high-precision results. The implementation of more generally appropriate generators is discussed in Ref. [12].

For the benchmarks, we compared the performance of the outlined GPU implementation on a Tesla C1060 as well as a more recent GTX 480 of the Fermi architecture series with the results of an optimized, single-threaded CPU code running on an Intel Core 2 Quad Q9650 at 3.0 GHz. For the Ising ferromagnet, $n = 1$ and $J_{ij} = 1$ in Eq. (1), we find a tile size of 16×16 spins to be optimal for sufficiently large systems [12]. The maximum performance reached on the GTX 480 is around 0.03 ns per spin flip (using $k = 100$), which is 235 times faster than the CPU implementation, cf. the data collected in Tab. 1 and the right panel of Fig. 1. The Tesla C1060, on the other hand, roughly performs at half of this speed. This speed-up, however, is only reached for sufficiently large system sizes that allow to fully load the 240 and 480 cores of the C1060 and the GTX 480, respectively. Very similar relative performance is observed for the Ising model in three dimensions. For models with continuous spins, exemplified by the 2D Heisenberg model with $n = 3$ and $J_{ij} = 1$ in Eq. (1), issues of floating-point performance and precision become important. We find that a mixed-precision calculation, where the spins are represented in single precision and only aggregate quantities such as the total energy are calculated in double precision (see “Metropolis single” in Tab. 1) yield high performance without problems with precision. If we employ the hardware optimized implementations of the special functions (trigonometric, exponential, logarithmic etc.) provided in the CUDA framework total speed-ups beyond 1000 can be achieved as compared to CPU codes (see “Metropolis fast math” in Tab. 1).

Simulations of systems with quenched disorder allow for trivial parallelization over disorder realizations on top of the domain decomposition outlined above. For the Edwards-Anderson Ising spin glass with couplings $J_{ij} \in \{-J, J\}$




Figure 1: Left: double checkerboard decomposition of a square lattice of edge length $L = 32$ for performing a single spin-flip Metropolis simulation of spin models on GPU. Right: Time per spin-flip in the 2D Ising model on CPU and on GPU with different choices of k . GPU data are for the Tesla C1060 device apart from the lowest curve which is for a GTX 480 card.

drawn from a bimodal distribution, we again find speed-ups of around 100 and 200 for the Tesla C1060 and the GTX 480, respectively. Further improvements can be achieved on using 64-bit multi-spin coding which allows for spin-flip times down to 2 pico-seconds on the GTX 480, cf. Tab. 1.

3. Cluster updates

While single spin-flip simulations on a fixed lattice appear to be near optimal problems for the parallel compute model of GPUs, highly non-local updates such as the cluster algorithms used to beat critical slowing down in ferromagnetic models close to a continuous phase transition are significantly harder to efficiently implement in parallel. To test this, we considered different implementations of the Swendsen-Wang cluster algorithm [14] for the Ising model. An update consists of the following steps:

1. Activate bonds between like spins with probability $p = 1 - e^{-2\beta J}$.
2. Construct (Swendsen-Wang) spin clusters from domains connected by active bonds.
3. Flip independent clusters with probability $1/2$.

While steps 1 and 3 are completely local and hence can be easily performed in a highly parallel fashion using a single thread for updating a few bonds or sites, the cluster identification step is intrinsically non-local, in particular close to the critical point where the Swendsen-Wang clusters undergo a percolation transition. The approach we adopt uses two steps: (a) identify clusters in tiles of the system, disregarding any couplings crossing the tile boundaries and (b) amalgamate those small clusters by successively taking the boundary bonds into account. From the approaches for clustering in tiles we implemented, including the Hoshen and Kopelman algorithm, breadth-first search, union-and-find algorithms [15] and “self-labeling” [16], we find the latter to be most efficient for intermediate tile sizes. For self-labeling each site starts out with a unique cluster label. Subsequently, each site in parallel sets its cluster label to the minimum of its own label and that of its northward and eastward neighbors (in two dimensions). This procedure is iterated until no change in cluster labels occurs. A maximum of the order of T iterations is required for tiles of size T , such that in total T^3 operations are necessary. This appears unfavorable compared to the other algorithms for local cluster labeling, but this disadvantage is counter-balanced for tile sizes up to about $T = 16$ by the high degree of parallelism. From a number of different options for the consolidation of cluster labels between tiles, including an iterative relaxation procedure and a hierarchical approach based on union-and-find data structures, we find the latter to




Figure 2: Left: speed-up of a Swendsen-Wang cluster update simulation of the 2D Ising model on a GTX 480 GPU as compared to the reference implementation on an Intel Q9650 CPU. Right: speed-up of the windowed multicanonical (MUCA) and Wang-Landau (WL) simulations of the 2D Ising model on the GTX 480.

be significantly more efficient. As is seen from the left panel of Fig. 2, this allows for up to around 20-fold speed-ups as compared to the CPU reference code, which appears respectable in view of the lack of locality, but is clearly less impressive than the results for local simulations.

4. Multicanonical and Wang-Landau simulations

Simulations of systems undergoing first-order phase transitions or featuring complex free-energy landscapes suggest the use of generalized-ensemble methods such as multicanonical (MUCA) [17] or Wang-Landau (WL) [18] simulations. Considering the internal energy E as reaction coordinate, the canonical distribution $p_\beta(E) = Z_\beta^{-1} \Omega(E) \exp(-\beta E)$ is generalized to read $p_\beta(E) = Z_\beta^{-1} \Omega(E) / W(E)$. A flat histogram is reached if the weights $W(E)$ equal the density of states $\Omega(E)$ or, equivalently, $\omega(E) \equiv \ln W(E) = S(E)$, where $S(E)$ is the microcanonical entropy. While MUCA uses a series of fixed-ensemble, equilibrium simulations to estimate $W(E) = S(E)$, an analogous estimate is calculated online in a non-equilibrium simulation in the WL approach. These algorithms are difficult to parallelize since they require knowledge of the current value of a global reaction coordinate (such as energy or magnetization) prior to each update. This effectively serializes all updates performed on a single instance of the system. To still benefit from the parallel GPU architecture, we use “windowing”, i.e., the idea of applying algorithms separately in small, fixed energy windows [18] and gluing together the resulting estimates to reconstruct the overall $S(E)$. We also use trivial parallelization to improve statistics and estimate statistical errors.

We implemented MUCA and WL codes on GPU, where each thread works on a separate copy of the system in a fixed individual energy window. Adjacent windows overlap by one energy value. To construct an initial spin configuration inside of the desired window we randomly select one among two nearest (AF or FM) ground states, and change its energy by the steepest ascent method for randomly selected spins. Each block of threads works on the same energy window to minimize thread divergence. Energy histograms and $S(E)$ estimates are stored in shared memory. Spin configurations are either directly worked on in global memory or tiled in shared memory. For a sufficiently high load of the GPU and ensuring coalescence of memory accesses, we do not find any benefit of using shared memory as each spin is only updated once before it is written back to global memory. All calculations are done in single precision, leading to essentially identical results as the double-precision CPU implementation.

For the 2D Ising model, we find that “windowing” does not cause systematic deviations from the exact result for $S(E)$ [19] as long as enough statistics is collected in each window. For the WL algorithm this means imposing a strict

System	Algorithm	L	CPU	C1060	GTX 480	speed-up
			ns/flip	ns/flip	ns/flip	
2D Ising	Metropolis	32	8.3	2.58	1.60	3/5
2D Ising	Metropolis	16 384	8.0	0.077	0.034	103/235
3D Ising	Metropolis	512	14.0	0.13	0.067	107/209
2D Heisenberg	Metro. double	4096	183.7	4.66	1.94	39/95
2D Heisenberg	Metro. single	4096	183.2	0.74	0.50	248/366
2D Heisenberg	Metro. fast math	4096	183.2	0.30	0.18	611/1018
2D spin glass	Metropolis	32	14.6	0.15	0.070	97/209
2D spin glass	Metro. multi-spin	32	0.18	0.0075	0.0023	24/78
2D Ising	Cluster	8192	77.4	—	4.62	—/17
2D Ising	multicanonical	64	42.1	—	0.33	—/128
2D Ising	Wang-Landau	64	43.6	—	0.94	—/46

Table 1: Spin-flip times for simulations of various lattice spin models with different algorithms on an Intel Q9650, a Tesla C1060 and a GTX 480, respectively. Apart from the cluster, multicanonical and Wang-Landau simulations, multi-hit updates with $k = 100$ have been employed.

criterion as to when the energy histogram is considered flat; for the MUCA algorithm a sufficient number of tunneling events should be demanded. This allows us, for instance, to construct $S(E)$ for a 64×64 system from windows as small as $\Delta E = 16$. The speedups of the GPU implementations using a sufficiently large number of windows and independent runs to fully load the GPU are depicted in the right panel of Fig. 2 and summarized in Table 1. For MUCA, we arrive at a speedup of 128, similar to the results found for the local algorithms, whereas the WL approach, in its current implementation, allows a 46 times performance increase only. This difference results from the dynamical nature of the WL algorithm, where run times are random variables, which leads to thread divergence and idle cores on the GPU. A more sophisticated implementation using some load balancing scheme will be discussed elsewhere.

Acknowledgments

The authors acknowledge support by the “Center for Computational Sciences in Mainz” (SRFN). M.W. received support by the DFG through the Emmy Noether Programme under contract No. WE4425/1-1.

- [1] K. Binder, D. P. Landau, A Guide to Monte Carlo Simulations in Statistical Physics, 2nd edition, Cambridge University Press, 2005.
- [2] H. W. J. Blöte, L. N. Shchur, A. L. Talapov, The cluster processor: New results, Int. J. Mod. Phys. C 10 (1999) 1137.
- [3] F. Belletti *et al.*, Janus: An FPGA-based system for high-performance scientific computing, Comput. Sci. Eng. 11 (2009) 48–58.
- [4] J. D. Owens, M. Houston, D. Luebke, S. Green, J. E. Stone, J. C. Phillips, GPU computing, Proc. IEEE 96 (2008) 879–899.
- [5] W. W. Hwu (Ed.), GPU Computing Gems: Emerald Edition, Morgan Kaufmann, Amsterdam, 2011.
- [6] D. B. Kirk, W. W. Hwu, Programming Massively Parallel Processors, Elsevier, Amsterdam, 2010.
- [7] OpenCL, <http://www.khronos.org/opencl/>.
- [8] CUDA zone, http://www.nvidia.com/object/cuda_home_new.html.
- [9] N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, E. Teller, Equation of state calculations by fast computing machines, J. Chem. Phys. 21 (1953) 1087.
- [10] D. Heermann, Parallelization of the Ising model and its performance evaluation, Parallel Comput. 13 (3) (1990) 345–357.
- [11] M. Weigel, Performance potential for simulating spin models on GPU, Preprint arXiv:1101.1427.
- [12] M. Weigel, Simulating spin models on GPU, Comput. Phys. Commun. 182 (2011) 1833.
- [13] <http://www.cond-mat.physik.uni-mainz.de/~weigel/GPU>.
- [14] R. H. Swendsen, J. S. Wang, Nonuniversal critical dynamics in Monte Carlo simulations, Phys. Rev. Lett. 58 (1987) 86.
- [15] M. E. J. Newman, R. M. Ziff, Fast monte carlo algorithm for site or bond percolation, Phys. Rev. E 64 (2001) 016706.
- [16] C. F. Baillie, P. D. Coddington, Cluster identification algorithms for spin models sequential and parallel, Concurrency: Pract. Exper. 3 (1991) 129–144.
- [17] B. A. Berg and T. Neuhaus, Multicanonical ensemble: A new approach to simulate first-order phase transitions, Phys. Rev. Lett. 68 (1992) 9.
- [18] F. Wang and D. P. Landau, Determining the density of states for classical statistical models: A random walk algorithm to produce a flat histogram, Phys. Rev. E 64 (2001) 056101.
- [19] P. D. Beale, Exact distribution of energies in the two-dimensional Ising model, Phys. Rev. Lett. 76 (1996) 78.